

In [1]:

```
# Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
```

Using TensorFlow backend.

from **future** import print_function read about it <https://www.pythonsheets.com/notes/python-future.html>

In [2]:

```
batch_size = 128
num_classes = 10
epochs = 10

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if K.image_data_format() == 'channels_first':
    #https://machinelearningmastery.com/a-gentle-introduction-to-channels-first-and-channels-last-image-formats-for-deep-learning/ to know about image_data_format and what is "channelS_first"
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 2s 0us/step

In [3]:

```
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255 #normalizing
x_test /= 255 #normalizing
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

In [4]:

```
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/11490434
```

```
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

Category 1: Kernel_size (3,3) without adding dropout and batch normalization

2 Hidden Layers

In [5]:

```
model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Conv2D(64, (3, 3), activation='relu'))
#model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
conv2d_2 (Conv2D)	(None, 24, 24, 128)	36992
conv2d_3 (Conv2D)	(None, 22, 22, 64)	73792
flatten_1 (Flatten)	(None, 30976)	0
dense_1 (Dense)	(None, 10)	309770
Total params: 420,874		
Trainable params: 420,874		
Non-trainable params: 0		

In [6]:

```
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 533s 9ms/step - loss: 0.1714 - accuracy: 0.9485 - v
al_loss: 0.0578 - val_accuracy: 0.9804
Epoch 2/10
60000/60000 [=====] - 534s 9ms/step - loss: 0.0509 - accuracy: 0.9844 - v
al_loss: 0.0400 - val_accuracy: 0.9867
Epoch 3/10
60000/60000 [=====] - 531s 9ms/step - loss: 0.0352 - accuracy: 0.9892 - v
al_loss: 0.0351 - val_accuracy: 0.9882
Epoch 4/10
```

```

60000/60000 [=====] - 529s 9ms/step - loss: 0.0256 - accuracy: 0.9925 - v
al_loss: 0.0306 - val_accuracy: 0.9904
Epoch 5/10
60000/60000 [=====] - 530s 9ms/step - loss: 0.0192 - accuracy: 0.9942 - v
al_loss: 0.0331 - val_accuracy: 0.9892
Epoch 6/10
60000/60000 [=====] - 530s 9ms/step - loss: 0.0146 - accuracy: 0.9957 - v
al_loss: 0.0366 - val_accuracy: 0.9890
Epoch 7/10
60000/60000 [=====] - 537s 9ms/step - loss: 0.0111 - accuracy: 0.9969 - v
al_loss: 0.0375 - val_accuracy: 0.9890
Epoch 8/10
60000/60000 [=====] - 537s 9ms/step - loss: 0.0090 - accuracy: 0.9975 - v
al_loss: 0.0382 - val_accuracy: 0.9893
Epoch 9/10
60000/60000 [=====] - 538s 9ms/step - loss: 0.0066 - accuracy: 0.9981 - v
al_loss: 0.0408 - val_accuracy: 0.9892
Epoch 10/10
60000/60000 [=====] - 534s 9ms/step - loss: 0.0049 - accuracy: 0.9986 - v
al_loss: 0.0450 - val_accuracy: 0.9890
Test loss: 0.04503656446120733
Test accuracy: 0.9890000224113464

```


In [7]:

```

import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.04503656446120733
Test accuracy: 0.9890000224113464



3 Hidden Layers

In [8]:

```
model = Sequential()
```

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 26, 26, 32)	320
conv2d_5 (Conv2D)	(None, 24, 24, 128)	36992
conv2d_6 (Conv2D)	(None, 22, 22, 64)	73792
conv2d_7 (Conv2D)	(None, 20, 20, 32)	18464
flatten_2 (Flatten)	(None, 12800)	0
dense_2 (Dense)	(None, 10)	128010
Total params: 257,578		
Trainable params: 257,578		
Non-trainable params: 0		

In [9]:

```

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/10
60000/60000 [=====] - 613s 10ms/step - loss: 0.1503 - accuracy: 0.9539 -
val_loss: 0.0659 - val_accuracy: 0.9804
Epoch 2/10
60000/60000 [=====] - 615s 10ms/step - loss: 0.0483 - accuracy: 0.9854 -
val_loss: 0.0517 - val_accuracy: 0.9855
Epoch 3/10
60000/60000 [=====] - 620s 10ms/step - loss: 0.0322 - accuracy: 0.9898 -
val_loss: 0.0348 - val_accuracy: 0.9880
Epoch 4/10
60000/60000 [=====] - 621s 10ms/step - loss: 0.0245 - accuracy: 0.9927 -
val_loss: 0.0373 - val_accuracy: 0.9884
Epoch 5/10
60000/60000 [=====] - 620s 10ms/step - loss: 0.0180 - accuracy: 0.9942 -
val_loss: 0.0380 - val_accuracy: 0.9893
Epoch 6/10
60000/60000 [=====] - 613s 10ms/step - loss: 0.0149 - accuracy: 0.9949 -
val_loss: 0.0383 - val_accuracy: 0.9887
Epoch 7/10
60000/60000 [=====] - 614s 10ms/step - loss: 0.0114 - accuracy: 0.9963 -
val_loss: 0.0488 - val_accuracy: 0.9870
Epoch 8/10
60000/60000 [=====] - 617s 10ms/step - loss: 0.0092 - accuracy: 0.9969 -
val_loss: 0.0576 - val_accuracy: 0.9853
Epoch 9/10
60000/60000 [=====] - 625s 10ms/step - loss: 0.0078 - accuracy: 0.9974 -


```

```
val_loss: 0.0549 - val_accuracy: 0.9882
Epoch 10/10
60000/60000 [=====] - 626s 10ms/step - loss: 0.0067 - accuracy: 0.9979 -
val_loss: 0.0513 - val_accuracy: 0.9886
Test loss: 0.05131749645890795
Test accuracy: 0.9886000156402588
```

In [10]:

```
import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.05131749645890795
Test accuracy: 0.9886000156402588
```



Category 2: Kernel_size (5,5) , max_pooling with adding dropout and batch normalization

2Hidden Layer

In [11]:

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

model = Sequential()
model.add(Conv2D(32, kernel_size=(5,5),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Conv2D(64, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
```

```

model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 24, 24, 32)	832
conv2d_9 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_1 (MaxPooling2D)	(None, 10, 10, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 10, 10, 64)	256
dropout_1 (Dropout)	(None, 10, 10, 64)	0
conv2d_10 (Conv2D)	(None, 6, 6, 64)	102464
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 64)	0
batch_normalization_2 (Batch Normalization)	(None, 3, 3, 64)	256
dropout_2 (Dropout)	(None, 3, 3, 64)	0
flatten_3 (Flatten)	(None, 576)	0
dense_3 (Dense)	(None, 10)	5770
Total params: 160,842		
Trainable params: 160,586		
Non-trainable params: 256		

In [12]:

```

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/10
60000/60000 [=====] - 292s 5ms/step - loss: 0.2171 - accuracy: 0.9330 - val_loss: 0.1368 - val_accuracy: 0.9620
Epoch 2/10
60000/60000 [=====] - 283s 5ms/step - loss: 0.0770 - accuracy: 0.9762 - val_loss: 0.0345 - val_accuracy: 0.9896
Epoch 3/10
60000/60000 [=====] - 288s 5ms/step - loss: 0.0571 - accuracy: 0.9826 - val_loss: 0.0266 - val_accuracy: 0.9901
Epoch 4/10
60000/60000 [=====] - 284s 5ms/step - loss: 0.0479 - accuracy: 0.9850 - val_loss: 0.0257 - val_accuracy: 0.9905
Epoch 5/10
60000/60000 [=====] - 290s 5ms/step - loss: 0.0418 - accuracy: 0.9874 - val_loss: 0.0261 - val_accuracy: 0.9911
Epoch 6/10
60000/60000 [=====] - 284s 5ms/step - loss: 0.0368 - accuracy: 0.9881 - val_loss: 0.0251 - val_accuracy: 0.9916
Epoch 7/10
60000/60000 [=====] - 288s 5ms/step - loss: 0.0347 - accuracy: 0.9889 - v

```

```

al_loss: 0.0279 - val_accuracy: 0.9897
Epoch 8/10
60000/60000 [=====] - 284s 5ms/step - loss: 0.0299 - accuracy: 0.9905 - v
al_loss: 0.0220 - val_accuracy: 0.9932
Epoch 9/10
60000/60000 [=====] - 286s 5ms/step - loss: 0.0289 - accuracy: 0.9912 - v
al_loss: 0.0202 - val_accuracy: 0.9930
Epoch 10/10
60000/60000 [=====] - 279s 5ms/step - loss: 0.0255 - accuracy: 0.9918 - v
al_loss: 0.0179 - val_accuracy: 0.9938
Test loss: 0.017866041018166287
Test accuracy: 0.9937999844551086

```

In [14]:

```


import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.017866041018166287
Test accuracy: 0.9937999844551086

```



Category 3: Kernel_size (2,2) and strides=(2,2),max_pooling, padding = "same" with adding dropout and batch normalization

In [15]:

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(2,2),
                strides=(2,2),
                padding='same',
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (2, 2), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

```

```

model.add(Conv2D(64, (2, 2), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 14, 14, 32)	160
conv2d_12 (Conv2D)	(None, 13, 13, 64)	8256
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 64)	0
batch_normalization_3 (Batch Normalization)	(None, 6, 6, 64)	256
dropout_3 (Dropout)	(None, 6, 6, 64)	0
conv2d_13 (Conv2D)	(None, 5, 5, 64)	16448
max_pooling2d_4 (MaxPooling2D)	(None, 2, 2, 64)	0
batch_normalization_4 (Batch Normalization)	(None, 2, 2, 64)	256
dropout_4 (Dropout)	(None, 2, 2, 64)	0
flatten_4 (Flatten)	(None, 256)	0
dense_4 (Dense)	(None, 10)	2570
Total params: 27,946		
Trainable params: 27,690		
Non-trainable params: 256		

In [16]:

```

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 35s 591us/step - loss: 0.6308 - accuracy: 0.8009 -
val_loss: 0.3981 - val_accuracy: 0.9075
Epoch 2/10
60000/60000 [=====] - 36s 596us/step - loss: 0.2360 - accuracy: 0.9259 -
val_loss: 0.0935 - val_accuracy: 0.9713
Epoch 3/10
60000/60000 [=====] - 36s 598us/step - loss: 0.1894 - accuracy: 0.9406 -
val_loss: 0.0794 - val_accuracy: 0.9767
Epoch 4/10
60000/60000 [=====] - 37s 613us/step - loss: 0.1644 - accuracy: 0.9488 -
val_loss: 0.0733 - val_accuracy: 0.9778
Epoch 5/10
60000/60000 [=====] - 36s 592us/step - loss: 0.1497 - accuracy: 0.9529 -
val_loss: 0.0650 - val_accuracy: 0.9792
Epoch 6/10
60000/60000 [=====] - 36s 596us/step - loss: 0.1416 - accuracy: 0.9562 -

```



```

00000/00000 [=====] - 36s 604us/step - loss: 0.1222 - accuracy: 0.9616 -
val_loss: 0.0707 - val_accuracy: 0.9769
Epoch 7/10
60000/60000 [=====] - 36s 605us/step - loss: 0.1334 - accuracy: 0.9578 -
val_loss: 0.0609 - val_accuracy: 0.9814
Epoch 8/10
60000/60000 [=====] - 36s 601us/step - loss: 0.1281 - accuracy: 0.9592 -
val_loss: 0.0586 - val_accuracy: 0.9819
Epoch 9/10
60000/60000 [=====] - 36s 603us/step - loss: 0.1261 - accuracy: 0.9608 -
val_loss: 0.0568 - val_accuracy: 0.9823
Epoch 10/10
60000/60000 [=====] - 36s 604us/step - loss: 0.1222 - accuracy: 0.9616 -
val_loss: 0.0532 - val_accuracy: 0.9839
Test loss: 0.05317528839812148
Test accuracy: 0.9839000105857849

```


In [18]:

```

import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
# lvalidation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.05317528839812148
Test accuracy: 0.9839000105857849



Category 4: Kernel_size (7,7),max_pooling, padding = "valid" with adding dropout and batch normalization

In [21]:

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(7,7),
                padding='valid',
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(128, (7, 7), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

```

```

model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Conv2D(64, (7, 7), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()

```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
conv2d_20 (Conv2D)	(None, 22, 22, 32)	1600
conv2d_21 (Conv2D)	(None, 16, 16, 128)	200832
max_pooling2d_7 (MaxPooling2D)	(None, 8, 8, 128)	0
batch_normalization_7 (Batch Normalization)	(None, 8, 8, 128)	512
dropout_7 (Dropout)	(None, 8, 8, 128)	0
conv2d_22 (Conv2D)	(None, 2, 2, 64)	401472
max_pooling2d_8 (MaxPooling2D)	(None, 1, 1, 64)	0
batch_normalization_8 (Batch Normalization)	(None, 1, 1, 64)	256
dropout_8 (Dropout)	(None, 1, 1, 64)	0
flatten_5 (Flatten)	(None, 64)	0
dense_5 (Dense)	(None, 10)	650
=====		
Total params: 605,322		
Trainable params: 604,938		
Non-trainable params: 384		

In [22]:

```

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 510s 9ms/step - loss: 0.1752 - accuracy: 0.9494 - val_loss: 0.0913 - val_accuracy: 0.9845
Epoch 2/10
60000/60000 [=====] - 508s 8ms/step - loss: 0.0643 - accuracy: 0.9811 - val_loss: 0.0288 - val_accuracy: 0.9901
Epoch 3/10
60000/60000 [=====] - 510s 8ms/step - loss: 0.0507 - accuracy: 0.9851 - val_loss: 0.0322 - val_accuracy: 0.9895
Epoch 4/10
60000/60000 [=====] - 514s 9ms/step - loss: 0.0423 - accuracy: 0.9875 - val_loss: 0.0214 - val_accuracy: 0.9931
Epoch 5/10
60000/60000 [=====] - 515s 9ms/step - loss: 0.0356 - accuracy: 0.9891 - val_loss: 0.0114 - val_accuracy: 0.9951

```

```

al_loss: 0.0175 - val_accuracy: 0.9945
Epoch 6/10
60000/60000 [=====] - 514s 9ms/step - loss: 0.0320 - accuracy: 0.9904 - v
al_loss: 0.0218 - val_accuracy: 0.9922
Epoch 7/10
60000/60000 [=====] - 510s 8ms/step - loss: 0.0284 - accuracy: 0.9914 - v
al_loss: 0.0211 - val_accuracy: 0.9923
Epoch 8/10
60000/60000 [=====] - 514s 9ms/step - loss: 0.0285 - accuracy: 0.9915 - v
al_loss: 0.0309 - val_accuracy: 0.9909
Epoch 9/10
60000/60000 [=====] - 506s 8ms/step - loss: 0.0259 - accuracy: 0.9922 - v
al_loss: 0.0170 - val_accuracy: 0.9946
Epoch 10/10
60000/60000 [=====] - 507s 8ms/step - loss: 0.0224 - accuracy: 0.9928 - v
al_loss: 0.0215 - val_accuracy: 0.9929
Test loss: 0.021513261749742378
Test accuracy: 0.992900013923645

```


In [23]:

```

import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.021513261749742378
Test accuracy: 0.992900013923645



Display Using PrettyTable

In [26]:

```

from prettytable import PrettyTable
x=PrettyTable()
x.field_names=["Model","#Hidden Layers","Kernel-
Size", "MaxPooling", "Dropout/BatchNormalization", "Optimizer", "Activation", "Accuracy"] #column
headers
x.add_row(["1.", "2 (128-64)", "3X3", "False", "False", "Adadelta", "ReLu", "0.989"])

```

```
x.add_row(["2.", "3(128-64-32)", "3X3", "False", "False", "Adam", "ReLu", "0.988"])
x.add_row(["3.", "2(64-64)", "5X5", "2X2", "True", "Adadelata", "ReLu", "0.993"])
x.add_row(["4.", "2(64-64)", "2X2(s=2,p='same')", "2X2", "True", "Adadelata", "ReLu", "0.983"])
x.add_row(["5.", "2(128-64)", "7X7(s=2,p='valid')", "2X2", "True", "Adam", "ReLu", "0.992"])
print(x)
```

Model	#Hidden Layers	Kernel-Size	MaxPooling	Dropout/BatchNormalization	Optimizer	Activation	Accuracy
1.	2(128-64)	3X3	False	False	Adadelata	ReLu	0.989
2.	3(128-64-32)	3X3	False	False	Adam	ReLu	0.988
3.	2(64-64)	5X5	2X2	True	Adadelata	ReLu	0.993
4.	2(64-64)	2X2(s=2,p='same')	2X2	True	Adadelata	ReLu	0.983
5.	2(128-64)	7X7(s=2,p='valid')	2X2	True	Adam	ReLu	0.992

In []: